

-----AI Engineering Delivery -----

1. “Tell me about yourself”

I have over 18 years of experience spanning software engineering, data, analytics and AI, with a significant part of my career in financial services.

In my current role at Aware Super, I lead the Advanced Analytics capability, which I established and operationalised. My role covers AI strategy, operating model, delivery framework, governance, and the productionisation and adoption of AI solutions across the organisation.

More recently, a major focus has been Generative and Agentic AI. I've designed Agentic workflows that combine enterprise data, AI reasoning and human-in-the-loop controls, as well as AI-driven decisioning solutions that integrate predictive models with Agentic AI to provide explainable recommendations.

Prior to Aware, I led Data Science at Thermo Fisher ANZ, where I delivered an enterprise knowledge solution using RAG and vector search that reduced average call handling time by around 40%.

As a next step I've increasingly focused not just on building AI models, but leverage AI across the PDLC, governance and knowledge foundations required to deliver AI safely at scale.

That's what attracted me to this CFS opportunity — taking that experience further and applying AI across the broader Product Delivery Lifecycle.

2. “Why are you interested in this role?”

What attracted me is the combination of AI, engineering and product delivery.

At Aware, I've established the framework for delivering **AI and Agentic AI solutions from discovery through to production, supported by governance and engineering practices.**

This role gives me the opportunity to take that experience further and **use AI to improve the broader product delivery lifecycle at enterprise scale.**

3. “Why are you looking to leave Aware Super?”

I've enjoyed my time at Aware and I'm proud of what we've built. I joined with the opportunity to establish and mature the Advanced Analytics capability, and we've made strong progress across AI strategy, delivery, governance and adoption.

I'm now looking for a role with a broader enterprise remit, where I can build on that experience and work more closely across AI, engineering and product delivery.

That's why this CFS role feels like a natural next step for me.

4. “What do you bring to this role?”

I'd probably highlight four things.

First, I bring a combination of **AI and engineering experience**. I started my career in software engineering and solution design and subsequently moved through data, analytics and AI leadership.

Second, I have experience **establishing AI capabilities**, rather than only delivering individual solutions. At Aware I've established the operating model, delivery framework and governance approach for Advanced Analytics.

Third, I have practical experience with **Generative and Agentic AI**, including RAG, vector search, multi-agent workflows, human-in-the-loop controls and AI decisioning.

And finally, I've spent a significant part of my career in **regulated financial services**, so I'm comfortable balancing innovation with architecture, security, risk, compliance and governance.

I think that combination allows me to bridge Product, Engineering, Data and AI rather than looking at AI as a standalone capability.

5. “What do you mean when you say you “established the AI delivery framework?”

I mean defining a consistent way of taking an AI opportunity from the business problem through discovery, solution design, development, testing, governance and ultimately production.

At Aware, that involved establishing how business, Data, Technology, Architecture and Risk work together, defining development and deployment practices, introducing CI/CD and automated testing, and embedding model validation and Responsible AI controls into the delivery lifecycle.

The objective was to move away from one-off, manually delivered models towards a repeatable way of delivering AI solutions safely and at scale.

One of the outcomes was reducing our AI/ML delivery cycle from around four months to under four weeks.

6. Tell me about an Agentic AI solution you've delivered

At Aware, one of the problems we were looking at was how to turn large volumes of unstructured member feedback and interactions into actionable CX insights.

I designed an Agentic AI architecture where interaction data such as surveys, Zendesk tickets and call transcripts is brought into Snowflake and then processed through specialised agents.

One agent analyses the verbatim information to identify sentiment, themes and pain points. Another looks at impact, priority and potential root cause. Rather than allowing the system to automatically act on everything, we introduced human-in-the-loop validation so the CX team can review and refine the findings before operational action.

We also designed evaluation and observability into the architecture — looking at groundedness, correctness, tracing, cost and token usage — together with security and Responsible AI controls.

The important thing for me was that it wasn't simply an LLM summarisation exercise. We were designing an end-to-end workflow where AI reasoning, enterprise data, human judgement and governance work together.

7. “Why did you use multiple agents rather than one LLM prompt”

I wouldn't use multiple agents simply because Agentic AI is fashionable. There needs to be a reason for decomposing the problem.

In this case, the tasks were different. Identifying sentiment, themes and pain points is different from assessing business impact, prioritisation and root cause.

Separating those responsibilities made the workflow easier to reason about, evaluate and govern. It also allows individual components to evolve without redesigning the whole solution.

My general principle is to use the simplest architecture that solves the problem. If one model call is sufficient, I wouldn't introduce multi-agent complexity.

“How do you evaluate an Agentic/Generative AI solution?”

I think evaluation has to be designed into the solution rather than added at the end.

There are generic measures such as groundedness, correctness and tool-selection accuracy, but enterprise use cases also need domain-specific evaluation.

For example, in the Aware CX solution we considered business-specific rubrics alongside standard LLM evaluation. I'd also look at operational measures such as latency, failures, agent/tool traces, token consumption and cost.

Then there's the human dimension. For higher-risk decisions, human validation and feedback provide another important evaluation mechanism.

Ultimately, technical evaluation needs to connect to the business outcome — whether the workflow is actually producing useful and trusted actions.

How do you approach Responsible AI and governance?

My approach is to embed governance into the delivery lifecycle rather than treat it as an approval immediately before production.

At Aware, I worked with Risk, Compliance and InfoSec to establish the AI/ML governance framework. That included model validation, Responsible AI considerations, lineage and risk controls.

For Generative and Agentic AI, I extend that thinking to areas such as PII protection, access controls, groundedness and hallucination risk, evaluation, observability and human oversight.

The goal is to create appropriate guardrails without making governance so manual that teams avoid using the framework.

10. Tell me about your AI decisioning / Next Best Action solution

At Aware, we're combining traditional predictive ML with Agentic AI rather than asking an LLM to make the underlying prediction.

The ML layer generates propensity scores and the important positive and negative drivers behind the prediction. The Agentic layer then takes those signals, combines them with member context and translates them into an understandable recommendation and next action for a frontline user.

That separation is deliberate. The predictive model does the numerical decisioning where it's strong, while the LLM is used for reasoning, context translation and natural-language explanation.

It also improves explainability because instead of showing someone a score like 0.82, we can explain why an action is being recommended and how the frontline person should respond.

An initial Zendesk PoC delivered around a 10x uplift in recommendation conversion.

Tell me about your experience with enterprise knowledge / RAG

Use Thermo Fisher, not Aware, as the primary example.

At Thermo Fisher, customer-service teams had to navigate complex product information across different knowledge sources, which increased call handling time.

We developed an enterprise knowledge solution using RAG and vector search. Product documents, FAQs, policies and knowledge-base content were ingested, parsed and chunked, with embeddings used to support retrieval.

We combined semantic/vector retrieval with metadata filtering so that the system retrieved relevant information within the appropriate business context. The Agentic layer then synthesised the retrieved information into a grounded response for the service representative.

An important design principle was grounding and traceability — the answer needed to come from enterprise knowledge rather than simply relying on what the LLM knew.

The solution reduced average call handling time by around 40%.

How would you enable AI across CFS's Product Delivery Lifecycle?

This is probably the most important CFS-specific question.

I wouldn't start with AI tools. I'd start by understanding the existing product delivery lifecycle and where the biggest friction points are.

I'd look across discovery and requirements, design, development, testing, deployment and operations and work with Product, Engineering and Architecture to identify where AI can materially improve productivity, quality or speed.

I'd then select a small number of high-value use cases and prove them through controlled pilots — perhaps knowledge discovery and requirements refinement in one area, or engineering and testing assistance in another.

In parallel, I'd establish the common foundations required to scale: secure enterprise knowledge, reusable AI services and patterns, integration with engineering tooling, evaluation and testing, observability, security and governance.

Once we've demonstrated value, I'd industrialise the successful patterns so individual teams don't create isolated AI solutions.

And I'd measure the outcome — adoption, cycle time, quality, reuse and engineering productivity — rather than measuring success by the number of AI tools deployed.

What experience do you have with AI-enabled PDL/SDLC?

Be precise here.

My strongest direct experience has been leading the end-to-end delivery of AI products and capabilities — from discovery and architecture through development, testing, governance, deployment and operationalisation.

I've also established many of the engineering foundations required to scale that delivery, including reusable patterns, CI/CD, automated testing, evaluation, governance and operational integration.

Where this CFS opportunity takes that further is applying AI itself more broadly across the traditional product delivery lifecycle — for example requirements, design, engineering and testing.

I understand those patterns and see a strong connection with the Agentic AI, enterprise knowledge and delivery capabilities I've already been building, but I wouldn't claim that I've already AI-enabled every stage of a traditional software-development lifecycle.

What experience do you have with knowledge graphs and ontologies?

Another answer where accuracy matters.

I haven't directly implemented an enterprise knowledge graph, so I wouldn't overstate that experience.

My closest practical experience has been building enterprise knowledge solutions using RAG, vector databases, metadata, taxonomies and structured business context.

At Thermo Fisher, for example, we built a knowledge solution that ingested enterprise content, applied metadata-aware retrieval and used that context to ground AI responses.

So I understand the problem that ontology and knowledge graphs are solving — creating consistent meaning, relationships and business context across enterprise information — and I would bring strong adjacent experience, but I'd work closely with enterprise information and architecture specialists where deeper ontology expertise is required.

How would you approach CFS's enterprise intelligence layer?

I'd think about it as more than a vector database.

The first layer is trusted enterprise information — identifying the authoritative knowledge and data sources.

Then you need the semantic structure around that information: metadata, business concepts, relationships, taxonomy and potentially ontology and knowledge graph capabilities.

Above that you need retrieval and knowledge services that AI applications and agents can consume consistently.

Then you need security, access control, evaluation, lineage and observability so those services can be trusted.

My experience at Thermo Fisher with enterprise RAG and at Aware with Agentic AI gives me experience particularly around the knowledge ingestion, retrieval, reasoning, governance and consumption layers. I'd complement that with the appropriate ontology and knowledge-graph expertise when designing the broader enterprise layer.

Build versus buy — how do you decide?

Your Aware architecture gives you a nice answer.

I start with the business requirement and architecture rather than assuming everything needs to be custom built.

I look at what the platform provides natively and where differentiation or enterprise requirements genuinely require custom engineering.

For example, with Snowflake Cortex, standard evaluation, observability, RBAC and some guardrails can be provided natively. But domain-specific evaluation, dynamic PII handling or tracing across external multi-agent orchestration may require custom components.

My preference is to use managed platform capability where it meets the requirement and invest engineering effort where it creates meaningful differentiation or addresses an enterprise control gap.

Tell me about a stakeholder conflict

At AMP, we developed an analytics solution that could reduce claims duration by around 20%. During development, Claims leadership changed and the new leadership team was sceptical about relying on the model.

Rather than pushing the solution, I spent time understanding their concerns and created a standalone testing environment where they could evaluate it without affecting BAU.

We then held regular sessions with Claims leaders and SMEs to review results and incorporate their feedback.

That transparency helped rebuild trust and we ultimately progressed to a production pilot.

The lesson for me was that successful AI adoption isn't primarily a technology problem. Trust, transparency and stakeholder ownership are equally important.

Tell me about a failure

Earlier in my leadership career, I led some analytics PoCs where I focused too heavily on proving the technical value before building sufficient stakeholder ownership.

The solutions performed well technically, but adoption was slower than I expected.

Looking back, I'd involved stakeholders in requirements, but not deeply enough in shaping the solution, defining success measures and planning the operational change.

Since then, I've changed my approach. I involve stakeholders earlier, agree success measures upfront, identify architecture, operational and risk impacts before development, and use smaller pilots to build confidence before scaling.

It changed how I think about AI delivery — a technically successful solution isn't successful if people don't adopt it.

How do you influence Engineering teams when they don't report to you?

I start with the problem we're trying to solve rather than trying to impose an AI solution.

Engineering teams need to understand why the capability matters, how it fits with existing architecture and tooling, and what impact it will have on reliability, security and support.

I involve Engineering and Architecture early in solution design, make the trade-offs transparent and establish clear ownership between the AI, platform and engineering components.

I've worked this way throughout my career because many of the teams I've needed to influence — Technology, Architecture, Risk, InfoSec and business teams — haven't reported directly to me.

What would your first 90 days look like?

Initially I'd focus on understanding rather than immediately changing things.

I'd understand CFS's current PDLC, AI strategy and architecture, engineering toolchain, knowledge landscape and the work already underway. I'd also spend time with Product, Engineering, Architecture, Data, Risk and Security leaders to understand their priorities and pain points.

From that I'd identify the highest-value opportunities for AI across the delivery lifecycle and assess the foundational gaps — particularly knowledge, tooling, governance and integration.

I'd then agree a small number of priority use cases and a roadmap that balances **quick proof of value with the longer-term enterprise foundations**.

By 90 days I'd want alignment on the target approach, success measures, priority pilots and the capabilities required to scale them.

----- Head of Data -----

“Tell me about yourself”

This will probably be the first and most important answer. Aim for 60–90 seconds, not your whole career history.

I have over 18 years of experience across data, analytics and AI, with a significant part of my career in financial services, including insurance, banking and superannuation.

In my current role at Aware Super, I lead the Advanced Analytics capability, which I established from the ground up. My role covers strategy, operating model, governance, platform capability and the delivery of analytics and AI solutions, working closely with Technology, Risk, Compliance and InfoSec.

Prior to that, I led the Data Science capability for Thermo Fisher ANZ, and before that I was Director of Data & Analytics at Cloudwick. **I also spent several years at AMP, including as Head of Customer & Market Analytics for the life insurance business, where I led a multidisciplinary team and was responsible for the data and analytics strategy and roadmap.**

So I've had the opportunity to work across the full data lifecycle—from strategy and governance through to platforms, engineering, analytics and AI—and what interests me now is taking that experience into a broader enterprise data leadership role.

Tell me about your experience with enterprise knowledge / RAG

At Thermo Fisher, customer-service teams had to navigate complex product information across different knowledge sources, which increased call handling time.

We developed an enterprise knowledge solution using RAG and vector search. Product documents, FAQs, policies and knowledge-base content were ingested, parsed and chunked, with embeddings used to support retrieval.

We combined semantic/vector retrieval with metadata filtering so that the system retrieved relevant information within the appropriate business context. The Agentic layer then synthesised the retrieved information into a grounded response for the service representative.

An important design principle was grounding and traceability — the answer needed to come from enterprise knowledge rather than simply relying on what the LLM knew.

The solution reduced average call handling time by around 40%.

How would you enable AI across CFS's Product Delivery Lifecycle?

I wouldn't start with AI tools. I'd start by understanding the existing product delivery lifecycle and where the biggest friction points are.

I'd look across discovery and requirements, design, development, testing, deployment and operations and work with Product, Engineering and Architecture to identify where AI can materially improve productivity, quality or speed.

I'd then select a small number of high-value use cases and prove them through controlled pilots — perhaps knowledge discovery and requirements refinement in one area, or engineering and testing assistance in another.

In parallel, I'd establish the common foundations required to scale: secure enterprise knowledge, reusable AI services and patterns, integration with engineering tooling, evaluation and testing, observability, security and governance.

Once we've demonstrated value, I'd industrialise the successful patterns so individual teams don't create isolated AI solutions.

And I'd measure the outcome — adoption, cycle time, quality, reuse and engineering productivity — rather than measuring success by the number of AI tools deployed.

What experience do you have with AI-enabled PDLC/SDLC?

My strongest direct experience has been leading the end-to-end delivery of AI products and capabilities — from discovery and architecture through development, testing, governance, deployment and operationalisation.

I've also established many of the engineering foundations required to scale that delivery, including reusable patterns, CI/CD, automated testing, evaluation, governance and operational integration.

Where this CFS opportunity takes that further is applying AI itself more broadly across the traditional product delivery lifecycle — for example requirements, design, engineering and testing.

I understand those patterns and see a strong connection with the Agentic AI, enterprise knowledge and delivery capabilities I've already been building, but I wouldn't claim that I've already AI-enabled every stage of a traditional software-development lifecycle.

What experience do you have with knowledge graphs and ontologies?

I haven't directly implemented an enterprise knowledge graph, so I wouldn't overstate that experience.

My closest practical experience has been building enterprise knowledge solutions using RAG, vector databases, metadata, taxonomies and structured business context.

At Thermo Fisher, for example, we built a knowledge solution that ingested enterprise content, applied metadata-aware retrieval and used that context to ground AI responses.

So I understand the problem that ontology and knowledge graphs are solving — creating consistent meaning, relationships and business context across enterprise information — and I would bring strong adjacent experience, but I'd work closely with enterprise information and architecture specialists where deeper ontology expertise is required.

How would you approach CFS's enterprise intelligence layer?

I'd think about it as more than a vector database.

The first layer is trusted enterprise information — identifying the authoritative knowledge and data sources.

Then you need the semantic structure around that information: metadata, business concepts, relationships, taxonomy and potentially ontology and knowledge graph capabilities.

Above that you need retrieval and knowledge services that AI applications and agents can consume consistently.

Then you need security, access control, evaluation, lineage and observability so those services can be trusted.

My experience at Thermo Fisher with enterprise RAG and at Aware with Agentic AI gives me experience particularly around the knowledge ingestion, retrieval, reasoning, governance and consumption layers. I'd complement that with the appropriate ontology and knowledge-graph expertise when designing the broader enterprise layer.

Build versus buy — how do you decide?

I start with the business requirement and architecture rather than assuming everything needs to be custom built.

I look at what the platform provides natively and where differentiation or enterprise requirements genuinely require custom engineering.

For example, with Snowflake Cortex, standard evaluation, observability, RBAC and some guardrails can be provided natively. But domain-specific evaluation, dynamic PII handling or tracing across external multi-agent orchestration may require custom components.

My preference is to use managed platform capability where it meets the requirement and invest engineering effort where it creates meaningful differentiation or addresses an enterprise control gap.

Tell me about a stakeholder conflict?

At AMP, we developed an analytics solution that could reduce claims duration by around 20%. During development, Claims leadership changed and the new leadership team was sceptical about relying on the model.

Rather than pushing the solution, I spent time understanding their concerns and created a standalone testing environment where they could evaluate it without affecting BAU.

We then held regular sessions with Claims leaders and SMEs to review results and incorporate their feedback.

That transparency helped rebuild trust and we ultimately progressed to a production pilot.

The lesson for me was that successful AI adoption isn't primarily a technology problem. Trust, transparency and stakeholder ownership are equally important.

Tell me about a failure

Earlier in my leadership career, I led some analytics PoCs where I focused too heavily on proving the technical value before building sufficient stakeholder ownership.

The solutions performed well technically, but adoption was slower than I expected.

Looking back, I'd involved stakeholders in requirements, but not deeply enough in shaping the solution, defining success measures and planning the operational change.

Since then, I've changed my approach. I involve stakeholders earlier, agree success measures upfront, identify architecture, operational and risk impacts before development, and use smaller pilots to build confidence before scaling.

It changed how I think about AI delivery — a technically successful solution isn't successful if people don't adopt it.

How do you influence Engineering teams when they don't report to you?

I start with the problem we're trying to solve rather than trying to impose an AI solution.

Engineering teams need to understand why the capability matters, how it fits with existing architecture and tooling, and what impact it will have on reliability, security and support.

I involve Engineering and Architecture early in solution design, make the trade-offs transparent and establish clear ownership between the AI, platform and engineering components.

I've worked this way throughout my career because many of the teams I've needed to influence — Technology, Architecture, Risk, InfoSec and business teams — haven't reported directly to me.

What would your first 90 days look like?

Initially I'd focus on understanding rather than immediately changing things.

I'd understand CFS's current PDLC, AI strategy and architecture, engineering toolchain, knowledge landscape and the work already underway. I'd also spend time with Product, Engineering, Architecture, Data, Risk and Security leaders to understand their priorities and pain points.

From that I'd identify the highest-value opportunities for AI across the delivery lifecycle and assess the foundational gaps — particularly knowledge, tooling, governance and integration.

I'd then agree a small number of priority use cases and a roadmap that balances quick proof of value with the longer-term enterprise foundations.

By 90 days I'd want alignment on the target approach, success measures, priority pilots and the capabilities required to scale them.

“Why are you interested in the Head of Data role?”

What attracted me is the breadth of the role. It's not just an analytics role—it brings together data strategy, governance, engineering, reporting and the enterprise data platform.

That's very closely aligned with how my career has evolved. I've developed data strategies and roadmaps, established governance frameworks, led multidisciplinary teams, to build impactful analytics capabilities. I have also worked on platform modernisation that has led improved efficiency and productivity .

The other attraction is insurance. I spent several years at AMP Life working across claims, underwriting, fraud and customer analytics, so I understand how important trusted data is in an insurance environment.

I also like that Adica is looking at how emerging technologies and AI can support the future data capability. That's an area where I can bring additional experience.

“Why are you looking to leave Aware Super?”

I've enjoyed my time at Aware and I'm proud of what we've built. I joined with an opportunity to establish and mature the Advanced Analytics capability, and we've made significant progress across strategy, governance, platform capability and delivery.

I'm now interested in taking the next step into a broader enterprise leadership position where I can take responsibility across the wider data capability—not only advanced analytics and AI, but also data strategy, governance, engineering, platforms and reporting.

That's why the Head of Data opportunity caught my attention.

“What do you bring to Adica?”

I think I bring three things that are particularly relevant.

First is breadth across data. I've worked across strategy, governance, platforms, engineering, BI, advanced analytics and AI, so I understand how those capabilities need to work together rather than operate as separate functions.

Second is financial services and insurance experience. At AMP I worked directly across life insurance functions, including claims, underwriting, fraud and customer analytics, and I've subsequently worked in superannuation, so I'm very comfortable operating in regulated environments.

Third is capability building. I've established and transformed data and analytics functions, built multidisciplinary teams and worked with senior business and technology stakeholders to move from strategy into actual delivery.

I also bring strong AI experience, which I see as complementary to the Head of Data role rather than separate from it.

“What are you looking for in your next role?”

I'm looking for a role where I can have broader enterprise impact and help shape the direction of the data capability.

I'm particularly interested in an environment where data is seen as a strategic business capability rather than purely a technology function, and where I can work closely with senior business and technology leaders to translate strategy into measurable outcomes.

Culture is also important to me. I enjoy building capable teams, giving people ownership and creating a collaborative environment where business, data and technology teams work together.

Your stakeholder-conflict / stakeholder-management story example

At AMP, we developed an analytics solution that could reduce claims duration by around 20%. During development, the Claims leadership changed and the new leadership team was sceptical about relying on the model.

Rather than pushing the solution, I captured their concerns and created a standalone testing environment where they could evaluate the model without affecting BAU. We then held regular sessions with Claims leaders and SMEs to review the results and incorporate their feedback.

That transparency helped rebuild trust and we ultimately progressed to a production pilot. The key learning for me was that stakeholder adoption is just as important as technical performance.

If asked about failure, take more ownership.

Earlier in my leadership career, I led some analytics PoCs where I focused too heavily on demonstrating the technical value of the solution before building sufficient stakeholder ownership around the problem.

The solutions performed well technically, but adoption was slower than I expected. I realised that I had involved stakeholders in defining requirements, but hadn't involved them deeply enough in shaping the solution and change journey.

Since then I've changed my approach. I involve business owners much earlier, agree success measures upfront, identify operational and risk impacts before development, and use smaller pilots to build confidence before scaling.

That experience changed how I lead transformation today.

How will I Lead Adica's Data Transformation

I would start by **understanding Adica's business strategy** and where data can create the greatest value, rather than starting with technology.

From there, I would **assess the current state** across the key areas — data platform and architecture, governance and data quality, reporting and analytics, people and capability, and how effectively data is being used across the business.

I would then work with the leadership team and key stakeholders across Technology, Actuarial and the business to **define a clear target state** and prioritised roadmap. I wouldn't try to transform everything at once. I'd identify a small number of high-value initiatives that can demonstrate value early, while progressively strengthening the underlying data foundations.

A key lesson from my experience is that data transformation isn't just a technology transformation. You can build a technically excellent platform or solution, but if the business doesn't trust it or adopt it, you haven't transformed

anything. So **I would involve business stakeholders early, establish clear ownership and success measures, and take them through the transformation journey.**

Finally, I would make sure governance, data quality, security and regulatory requirements are embedded into the way we deliver rather than treated as something we address at the end.

Ultimately, my objective would be to make data trusted, accessible and useful, so that it improves decision-making, operational efficiency and customer outcomes across Adica.

Introduction

I have over 18 years of experience in business consulting roles that have given me extensive exposure in providing solution recommendations to clients based on data and analytics. I bring strong expertise in developing data products and services using data science, data platforms and other advanced analytics techniques.

I have also led data & analytics teams and delivered solutions for a diverse set of client organisations ranging from Banking, Insurance, Super, Retail, Airline and Telecommunication sectors. I have successfully implemented data solutions for our clients, delivering operational efficiencies, enhanced customer experience, and minimised time to market. ~~In my recent role at Cloudwick Technologies, I was responsible for business development and leading client engagements.~~ In my current role at Aware, I was responsible for *leading the development of an impactful advanced analytics function to enhance data-driven decision-making to support member growth and retention..*

Escalator pitch - why I am interested in the Wespac role -

I see this role as a unique opportunity to lead impactful AI initiatives that deliver real business value. I'm passionate about applying data science in meaningful ways, and I've done this successfully across financial services — from building Next Best Action models and customer segmentation to AI solutions that improve operational efficiency and customer experience. What also excites me is the chance to drive innovation within Westpac's Business & Wealth division, while maintaining a strong focus on governance, risk, and compliance..

What do you bring to the company?

I have over 14 years of experience in business consulting roles that have given me extensive exposure in providing solution recommendations to clients based on data insights and business value. I also bring strong expertise in developing data products & services using data science and analytics platforms.

I have led advanced analytics teams and delivered solutions for a diverse set of client organisations ranging from Banking, Retail, Airline and Telecommunication sectors. This allows me to cross-pollinate ideas from different industries and innovate. In my current role at Cloudwick, I am responsible for business development and leading client engagements. I have successfully executed data solutions for our clients and delivered operational performance by integrating the final product.

Daiichi Sankyo  Daiichi is a global pharmaceutical company that is poised to redefine the treatment of cancer. Daiichi leverages science and technology to create innovative medicines for people with cancer, cardiovascular, and other diseases with high unmet medical needs.

Based on your experience, here's how you can confidently speak to **regulatory compliance, data governance, data quality, and ethical AI use** in a Lead Data Scientist interview—particularly relevant for a bank like **Westpac**, which operates in a highly regulated environment.

✓ 1. Regulatory Compliance

You can say:

“Working in superannuation and financial services, I’ve always ensured our models and data solutions comply with APRA regulations, GDPR (where applicable), and internal risk policies. For example, when developing a roll-out propensity model, I worked closely with **risk and compliance team** to ensure no discriminatory variables were used, and that we had full auditability on inputs and outputs—key requirements under responsible AI and model risk management standards.”

Key points you’ve likely experienced:

- Alignment with APRA’s CPG 235 (Data Risk Management).
 - Handling of Personally Identifiable Information (PII).
 - Documentation and model governance (e.g., version control, model validation).
-

✓ 2. Data Governance

You can say:

“Data governance was central to my work at AMP and in my current role. I collaborated with data stewards and business data owners to define clear data dictionaries and ensured lineage was maintained in Snowflake. We used role-based access controls and worked with Information Security teams to control sensitive datasets. This was crucial not just for compliance, but for building trust in our insights.”

Points you might emphasize:

- Use of **Snowflake** with **RBAC** and **data masking**.
 - Maintaining **data lineage** and documentation.
 - Collaborating with **data governance councils** or stewards.
-

✓ 3. Data Quality

You can say:

“No model is better than its data. That’s why I built robust data pipelines with quality checks embedded at every step—null checks, data drift monitoring, and reconciliation against source systems. For example, when building a member segmentation model, we validated inputs with SMEs to ensure consistency, and I implemented logging and alerts to catch anomalies in the data feed.”

Examples:

- Use of automated checks in Azure/Spark pipelines.
 - Validating business logic with SMEs.
 - Creating dashboards to monitor data quality KPIs.
-

✓ 4. Ethical Use of Data / AI

You can say:

“Ethical AI is a non-negotiable. I make it a priority to ensure our models avoid bias and maintain explainability. At AMP, I led the creation of a claim intervention model where we removed socioeconomic and demographic predictors that could unfairly impact vulnerable customers. I also used SHAP to provide explainability and worked with the business to ensure human-in-the-loop governance was in place.”

Key elements to touch on:

- Bias testing and fairness audits.
- Model explainability tools (e.g., SHAP, LIME).
- Avoiding black-box models in high-stakes use cases.

what motivates you?

I see this role as a good blend of leadership and expertise in data science to deliver business value to Daiichi in the ANZ region!

Based on my experience I have successfully led data analytics/science teams to support businesses by enabling data-driven decision-making, improving

operational efficiencies, and enhancing customer experience. I believe this is an opportunity that aligns well with my skill set!!

Questions -

What are the key challenges Superhero is looking to solve leveraging data & analytics to support its customers? I see in the job description, you have mentioned some key goals such as enabling self-service capabilities and democratising data organisation wide, but i am very keen to know what are your focus area for your customers.

What are the company's plans for growth and development based on it current business model?

Historically we have seen many Fintech startups fail, why do you think Superhero will be a success?

What do you want to find in a role and what are your expectations of the team and organisation?

1. An organization that is data-driven and able to provide me with challenges, where I can leverage my strong consulting and data science background to deliver value to the business.
2. A great set of people to work with, who respect each other
3. A career progression: the new position should be a setup in terms of influence and responsibilities

Kellogg's - A cereal behemoth and a global brand. I see this role as an opportunity to deliver impactful data solutions that align with Kellogg's business strategy for ANZ. This is something I have emulated very successfully across various organisations given my experience and expertise in leveraging data and analytics to solve complex business problems. I believe I am a very good fit for this role!

Culture Amp - Focuses on the culture of an organisation leveraging data & people science using advanced analytics techniques to identify critical insights to enhance employee engagement and performance

Pfizer - Pfizer as a leader in Healthcare industry with a major focus on leveraging Data Science for Drug Discovery, Identifying rare diseases, precision medicine and managing sensitive inventory & predict supply chain issues.

Why Data Ops. at WBC and what motivates you?

I see this role as a good blend of leadership and expertise in data science to deliver business value to 'Pfizer's' ANZ region!

Based on my experience I have successfully led data science teams to support businesses enabling data-driven decision-making, improving operational efficiencies, and enhancing customer experience. I believe this is an opportunity that aligns well with my skill set!!

I see this role as a good blend of leadership and expertise in data science to deliver business value to 'Culture Amp'.

Based on my experience I have successfully led data science teams to support businesses **enabling data-driven decision-making, improving operational efficiencies, and enhancing customer experience.** I believe this is an opportunity that aligns well with my skill set!!

Over 15 years of experience in business consulting roles with extensive exposure in providing strategic direction and solution recommendations to clients based on data and analytics.

I have led data & analytics teams while working at organisations across various industries. I have had the opportunity to set up and enhance data & analytics capabilities such as data management, governance and analytics tools. This has allowed me to deliver on strategic business goals for various organisations such as enabling data-driven decision making, improving operational efficiencies and enhancing the customer experience.

I believe this role aligns very well with my past experiences and skill-set.

I have over 15 years of experience in business consulting roles that have given me extensive exposure in providing solution recommendations to clients based on data and analytics. I bring strong expertise in developing data products and services using data platforms and advanced analytics techniques.

I see this as an opportunity that aligns well with my skill set. Over 15 years of experience in business consulting roles with extensive exposure in providing strategic direction and solution recommendations to clients based on data and analytics. .

I also have led data & analytics teams **and** while working at organisations across various industries. I have had the opportunity to setup and enhance data & analytics capabilities

to support businesses **enabling data-driven decision making, improving operational efficiencies and enhancing customer experience.** I believe this opportunity offers me a very similar challenge. But it also gives me an opportunity to work in a start-up fintech environment which to me is an exciting challenge in it self and I believe it will ample oppotnutiy to learn a lot from!

I see this as an opportunity that aligns well with my skill set. I bring over 15 years of experience in business consulting roles that have given me extensive exposure in providing solution recommendations to clients based on data and analytics.

What motivates me about this role is the challenge of exploring new opportunities to leverage data & analytics to recommend solutions to streamline business processes, and improve patient-centred health experience.

The role will give me ample opportunities to solve challenging business problems and recommend solutions to clients using advanced analytics techniques.

It also gives me the opportunity to try out new AA technologies and ideas, which is something I am very passionate about.

----- Sydney Water -----

I see this as a critical role playing important part in supporting the data and analytics department-wide. The role focuses on evolving the data & analytics capabilities going forward.

These challenges align very well with my past experiences, where I have successfully supported businesses by enabling data-driven decision making, delivering operational efficiencies leading to major cost savings and improving customer experience.

I see this role play an important part in supporting the data and analytics requirements for Standards Australia in their endeavour to develop and adopt quality standards to increase safety, sustainability and drive efficiencies across various industry sectors. The role also focuses on evolving the data & analytics capabilities going forward.

These challenges align very well with my past experiences, where I have successfully supported businesses by enabling data-driven decision making and delivering operational efficiencies and minimising time to market.

I see this role playing an important part in applying data science to innovate and leverage technologies to develop and deploy end-to-end AI solutions such as predictive analytics, computer vision, and NLP. This is an area I am very passionate about.

This also gives me the impression that the leadership have a clear vision of how technology could be used to solve some of the complex business problems. And from my experience evolving a data driven culture in an organisation cannot be achieved without the support of the C-level executive buy-in and support.

Annalect leverages data & analytics

Annalect **provides organisations with an environment and teams that democratize data insights and make it easy to apply them to media and creative executions.**

Data can help in -

- Personalising customer experience
- Help measure and improve campaign performance
- Optimising marketing ROI e.g. cost of conversion

How will you go about promoting AI in an organisation?

Refer to presentation...

https://docs.google.com/presentation/d/1kFgJpccCr3GZ7OjRvnVLh7n23myDnNidGSzJxgrv9IQ/edit#slide=id.g147b173b216_0_18

What is your work style and how do you engage with people?

I treat each team member fairly and equally. I am respectful, supportive, and empathetic to all my colleagues.

Leadership: I look at myself as a change agent with the ability to both influence and lead people. Let me give you a few examples of leadership from my recent past:

- In my current role at Thermo Fisher, I set up the Data Science function for ANZ and led the delivery of data solutions to accelerated business growth
- In my recent role at Cloudwick, I am heading the operations within the APAC region responsible for business development and leading client engagements. I successfully implemented data solutions for our clients delivering operational efficiencies and minimising time to market.
- At AMP, I supported the setup of advanced analytics functions within the company's life insurance group and delivered data solutions that led to significant cost savings and minimized risk impact on the business.
- Working for IBM, I recommended and supported the implementation of a cloud-based analytics platform for a leading investment bank.
- At CBA I pioneered the use of unstructured system logs to gain insights into user experiences and system performances.

Tell me about a time when you resolved a conflict with a stakeholder?

The business was looking to streamline the claims management process. The problem was put to the Analytics team to help solve. I designed and developed a data product that could reduce the claim duration by 20% and also help them make better allocation of resources within the Claims team. This was developed in close collaboration with the business SME's.

Unfortunately, at that point in time, the claims team went through a restructure and we lost the leadership support who helped us develop the data tool. With the induction of new claims team leadership, we were challenged to prove the value of the product to the business. In my presentation, I experience major resistance from the new claims team in terms of how dependable the tool would be. Even though I provided test results to them.

I took note of all the concerns that were raised by the claims team and came back with enhancements to the tool as prescribed by them. We created a standalone test environment to demonstrate the performance of the tool without interfering with the BAU process. I held fortnightly meetings with the new claims leaders and SME's to evaluate the performance of the tool. In a few sessions evaluating the performance of the tool with the claims team, we started getting positive feedback and were able to pilot the tool in production on death claims to start with.

The takeaway for me was a collaboration with the business team and transparency of the solutions help build confidence in the business, which are critical for the success of any data-driven initiative.

Tell me about a time you failed and what you have learned from it?

In my current role, I was given the mandate to leverage advanced analytics techniques to solve business problems. Discussing the business pain points with the business leadership team and I decided to come up with POC's in an effort to demonstrate how AI techniques could support the business in making informed decisions based on data. Some of my recommendations were to automate processes that would make them more efficient and result in cost savings.

But stakeholders were resistant to accept these solutions, due to high risk-averse culture. They also felt their assumptions were being challenged based on the data insights and test results.

This left the stakeholders uncomfortable and considered AI solutions as a threat to them. This stalled the uptake of the recommendations and solutions being proposed by the analytics team.

The learning I took away from this experience was being able to develop data tools that might be impactful doesn't guarantee success in leveraging analytics effectively in solving business problems. A data-driven culture is a key to the success of supporting an operating environment that seeks to leverage data whenever and wherever possible to enhance business efficiency and effectiveness. The most influential factor is the C-level executive buy-in and support, which frequently cascades down across to teams and the company to make it a data-driven culture.

Defining an AI strategy for an organization involves:

1. Align with Business Goals:

- Identify Key Objectives: Clearly define what problems AI can solve for your business.
- Prioritize Use Cases: Focus on areas with the highest potential impact.

2. Data Assessment:

- Data Availability: Evaluate the quality and accessibility of your data.
- Data Strategy: Develop a plan for data collection, storage, and management.

3. Talent and Skills:

- Skill Assessment: Identify necessary AI skills within your organization.
- Talent Acquisition: Plan for hiring or training AI talent.

4. Technology Infrastructure:

- Hardware and Software: Ensure you have the necessary computing resources.
- Cloud vs. On-Premise: Decide on the best deployment model.

5. Ethical Considerations:

- Bias and Fairness: Develop guidelines to address potential biases.
- Transparency: Plan for how to explain AI decisions.

6. Implementation and Governance:

- Pilot Projects: Start with smaller, focused AI projects.
- Governance Structure: Establish processes for decision-making and oversight.

7. Continuous Learning and Adaptation:

- Stay Updated: Monitor AI advancements and industry trends.
- Iterative Approach: Be prepared to adjust your strategy based on results.

Remember:

- **Involve Stakeholders:** Get input from different departments.
- **Set Realistic Expectations:** Understand that AI implementation takes time.
- **Measure Success:** Define key performance indicators (KPIs) to track progress.

How would you build an analytics team?

Building a strong analytics team requires a focus on both the technical aspects of data management and the team itself. Here's a breakdown of the key considerations:

Data:

- **Data collection capability:** Establish methods to gather relevant data from various sources.
- **Data quality and availability:** Ensure data accuracy, completeness, and accessibility for analysis.
- **ETL (Extract, Transform, Load):** Implement processes to move data from source systems to a central repository.

People:

- **Skill set:** Recruit individuals with data analysis, engineering, and communication expertise.
- **Professional development:** Provide opportunities for team members to learn new skills and stay updated with industry trends.
- **Morale:** Foster a collaborative and supportive work environment to keep team members engaged.

Infrastructure and tools:

- **Central data repository:** Create a secure and centralized location to store all your data.
- **Data sharing tools:** Facilitate easy access and sharing of data among authorized team members.
- **Data analysis tools:** Equip the team with appropriate software for data cleaning, analysis, and visualization.

Processes:

- **Delivery Framework (BADIR):** Define a framework for data analysis, such as Business Analysis, Data Acquisition, Data Integration, Reporting, and Data Interpretation.
- **Effective Communication:** Establish clear communication channels to ensure everyone is informed about project goals and progress.

- **Collaboration & Task management processes:** To effectively manage tasks and deadlines.

By addressing these aspects, you can build a well-rounded analytics team that can derive valuable insights from data to support informed decision-making.

a.

Do you manage people?

I take the following things into consideration in managing people:

1. **Foundational Framework:** Clear processes (which include the development framework, communication and feedback process), roles, and responsibilities provide a strong foundation for your team. This reduces ambiguity and provides direction for your team members.
2. **Team motivation:**
 - a. Align individual goals with team and company objectives to create a sense of purpose.
 - b. Workload management: Ensure workload is managed to acceptable levels.
 - c. Professional development (LEAP - Learn & Enable Action Plan): This ensures continuous improvement of team members and fosters creativity
 - d. Reward - Being suitably rewarded for their efforts
 - e. Empowering team members - manager should role-model mindsets and behavior that promote empowerment (e.g. knowledge sharing, helping each other). This also requires team significant support to get comfortable with failure
3. **Open Communication & Transparency:** Encourage open communication including Reciprocal feedback that allows for constructive feedback, transparency and space for information sharing. This builds trust and psychological safety within the team.
4. **Team building activities** and positive reinforcement

How to build a data-driven culture?

1. Data-driven culture starts at the (very) top
2. Fix basic data-access issues quickly - Data democratisation makes a huge impact
3. Identifying business metrics that matter is critical - It should also track the quality of those predictions – they will steadily improve!
4. Don't pigeonhole your data scientists.
5. Quantify uncertainty - **A test & learn culture** can help overcome uncertainty.
6. Make PoC's simple & quick. Keep it simple and robust, not fancy and brittle.
7. Specialised training should be provided JIT(Just in time)

8. Taking business stakeholders on a journey gives teams a deeper understanding of the approaches and often prompts them to consider a wider set of alternatives or to rethink fundamental assumptions

What do friends think of you?

Helpful, Resourceful & Trustworthy

What do your colleagues think of you?

Trustworthy, Supportive, Collaborative

What are your strengths?

- Ability to solve complex business problems leveraging data & analytics
- A change agent with the ability to both influence and lead people. I am flexible (ready for change, be creative and innovative to find non-traditional ways for a business to grow and exploit opportunities), a good listener, take responsibility & accountability, and work collaboratively with cross-functional teams.
- Strong stakeholder management skills.

Managing Conflicting Priorities

1. **Schedule Work Effectively** - Prioritising tasks is critical. You'll then know which tasks clash, which tasks you may need to negotiate new deadlines. This should be a collaborative effort with key stakeholders.
2. **Negotiate Deadlines** - Once you've clarified your priorities, look at your responsibilities and identify any deadlines that you need to re-negotiate.
3. **Managing Expectations** - It's also important that you manage stakeholder expectations before you take on new tasks.
4. **Being professional** - In times of stress especially during times of intense pressure, is what best demonstrates your character to those around you.

AMP Role


1. What would your vision for the team be under AMP Life - how to ensure we are working on the right things?

how we would cater to some of the other needs of the business (like some of the Customer analytics components which sit in the Customer Team and will eventually come to us), what success looks like?

2. How would you go about this?

3. Any thinking you have on how to develop the team and what the team would need to be successful in implementing an analytics function?

4. Any development areas you would see for yourself and the team to be successful.



Tackling bias in artificial intelligence (and in humans)

Minimizing bias will be critical if artificial intelligence is to reach its potential and increase people’s trust in the systems.

Six potential ways forward for artificial-intelligence (AI) practitioners and business and policy leaders to consider

1  Be aware of contexts in which AI can help correct for bias and those in which there is high risk for AI to exacerbate bias	2  Establish processes and practices to test for and mitigate bias in AI systems	3  Engage in fact-based conversations about potential biases in human decisions	4  Fully explore how humans and machines can best work together	5  Invest more in bias research, make more data available for research (while respecting privacy), and adopt a multidisciplinary approach	6  Invest more in diversifying the AI field itself
it is important to anticipate domains potentially prone to unfair bias	Operational strategies can include improving data collection through more cognizant sampling and using third parties to audit data and models.	Check if proxies used in the past are adequate and how AI can help by surfacing long-standing biases that may have gone unnoticed	Some promising systems use a combination of machines and humans to reduce bias. Techniques in this vein include “human-in-the-loop” decision making		A more diverse AI community will be better equipped to anticipate, spot, and review issues of unfair bias and better able to engage communities likely affected by bias

Unintended consequences of AI

While AI and advanced analytics offer many positive benefits, they can lead to significant unintended (or maliciously intended) consequences for individuals, organizations, and society.

Individuals

Physical safety



Privacy and reputation



Digital safety



Financial health



Equity and fair treatment



Organizations

Financial performance



Nonfinancial performance



Legal and compliance



Reputational integrity



Society

National security



Economic stability



Political stability



Infrastructure integrity



<https://www.mckinsey.com/business-functions/mckinsey-analytics/our-insights/confronting-the-risks-of-artificial-intelligence>

AI Governance

AI governance frameworks could help organizations learn, govern, monitor, and mature AI adoption. An AI governance framework might begin with an organization identifying key stakeholders representing various groups and departments. Such a 'coalition' may be formalized in a Center of Excellence (CoE), working group, or council, among other examples. Such groups might develop best practices for their organization, share knowledge, and build guardrails for the use of AI systems.

AI governance frameworks should also consider a host of other factors, some of which we outline below.

Monitoring and Oversight

A central monitoring and escalation process is key, where an organisation monitors the accuracy, consistency, and efficiency of decisions being made & processes being supported by AI.

Third-Party Risk Management

The use of AI/ML deployment may involve third party applications and/or data, as discussed in Section 2, which could enable scalability, increased compute power and access to vendors. The use of AI/ML deployment may involve third party applications and/or data, which could enable scalability, increased compute power and access to vendors.

Risk management may include testing certain aspects of current practices, such as TPRM transparency around model interpretability, information security issues for cloud-based service providers, and broader concerns around technology dependencies for the third parties themselves.

Roles and Responsibilities

Every organization is different with respect to their internal organizational structure and general roles and responsibilities. The roles/activities below provide some examples for organizations that are discussing roles and responsibilities with respect to AI implementations to consider.

Interpretability/Explainability

AI use cases where interpretability/explainability is required by law or is otherwise appropriate. AIRS (*Artificial Intelligence/Machine Learning (AI/ML) Risk and Security*) acknowledges that it may not always be necessary or appropriate. As of this writing, there is no commonly agreed upon standard definition of AI explainability.

Data Management

Data Governance - Data governance sets standards and policies for the **availability**, **usability**, **integrity**, and **security** of data. When done well, data governance helps you grow fast, and do so with consistency.

Data Governance vs. Data Management

While both are critical for effective data utilization, data governance and data management serve distinct purposes.

Data Governance

- **Focus:** Establishing policies, standards, and procedures for data management.
- **Role:** Sets the framework for how data should be handled, accessed, and used.
- **Activities:** Defining data ownership, data quality standards, security protocols, and compliance regulations.

Data Management

- **Focus:** Implementing the strategies and procedures outlined by data governance.
- **Role:** Handles the day-to-day operations of data, including collection, storage, processing, and analysis.
- **Activities:** Data cleaning, integration, storage optimization, and performance tuning.

To summarize:

- **Data governance** is about *planning* how data should be managed.
- **Data management** is about *doing* the work to manage data according to the plan.

Both are essential for ensuring data quality, security, and optimal use within an organization.